

Reflective Synthesis Paper

AI-Assisted GitOps Orchestration for Safe Helm Upgrades Across Multi-Cluster Kubernetes Platforms

Student: Sheunesu Gumbie

Programme: AI Mastery Capstone

1. Overview and Industry Context

Modern platform engineering teams are responsible for maintaining shared Kubernetes infrastructure across multiple departments, regions, and environments. A recurring operational challenge is the safe upgrade of platform components—such as Kafka Connect, Prometheus, Grafana, and Loki—which are installed through Helm charts and managed declaratively via GitOps tools such as Argo CD. Each upgrade cycle requires an engineer to: identify all affected clusters, research breaking changes, generate updated configuration, validate Helm charts, monitor INT deployments, and decide whether to promote changes to production. Across ten or more environments, this process is repetitive, error-prone, and inconsistently documented.

The information technology sector, and platform engineering specifically, is a domain where AI-assisted decision support can reduce toil and improve reliability without replacing the human judgment required for critical production systems. System-level interactions, data dependencies, and configuration complexity can create substantial hidden maintenance costs in AI-enabled systems (Sculley et al., 2015). The problem is not a lack of data—cluster inventory, upgrade history, release notes, and monitoring metrics are all available—but rather the cognitive overhead of synthesising them consistently across every upgrade cycle. This project addresses that problem with an integrated AI system that combines data analysis, machine learning-inspired risk scoring, generative AI synthesis, and agentic orchestration into a single, auditable upgrade workflow.

2. Integration Rationale

The system integrates methods from five prior capstone projects across four AI domains.

Project 2 — Statistical Data Analysis provided the EDA methodology that drives the inventory analysis phase. The same Pandas-based workflow used to profile a bank fraud dataset—completeness checks, `value_counts` distributions, groupby aggregations, and

descriptive statistics—is applied here to cluster inventory data to detect version drift, missing monitoring coverage, and configuration inconsistencies.

Project 3 — Applied Machine Learning (LSTM Fraud Detection) informed the risk scoring design. The LSTM classifier combined temporal features with threshold-based classification to distinguish fraud from legitimate activity. The upgrade risk model applies the same principle: upgrade features (version gap magnitude, breaking changes, Kubernetes compatibility, component criticality, historical success rate) are weighted and summed to a deterministic risk score mapping to a categorical risk level. The result is fully deterministic by design — a probabilistic gate on a production upgrade is unacceptable in a critical-infrastructure context.

Project 5 — Generative AI (VAE on Fashion-MNIST) demonstrated how a trained encoder compresses high-dimensional input into a compact, structured representation. This project applies a transferable design insight using a large language model: the LLM receives already-extracted, structured findings and synthesises them into a concise human-readable summary. Crucially, the LLM does not construct the safety-critical fields in the `ResearchReport`—those are built deterministically from tool observations. The LLM contributes only the synthesis note and optional INFO-level hypotheses.

Project 6a — Agentic AI Capstone (Tool-Constrained Research Assistant) contributed the bounded ReAct loop architecture that structures the `UpgradeResearchAgent` (Yao et al., 2023). The agent iterates up to five times, at each step producing a validated `ReActDecision`, invoking only tools registered in the `ToolRegistry`, and recording the result as a typed `ToolObservation`. The agent may not proceed to FINISH until all three mandatory evidence tools have succeeded. Tool arguments are scope-checked against the upgrade request, preventing the LLM from satisfying evidence requirements using data from a different component or version.

Project 6b — Agentic AI Beaver Choice (Multi-Agent Orchestration) contributed the sequential coordination pattern used in the `UpgradeCoordinator`. The workflow mirrors the paper supply quoting system's sequential stages, with structured inter-agent message passing through Pydantic models replacing raw text handoffs. Only the Research phase uses an LLM-assisted agent; the remaining stages are deterministic Python components — a deliberate choice to keep safety gates auditable and independent of LLM behaviour.

3. System Design and Technical Decisions

The system architecture separates two distinct layers: an AI reasoning layer and a deterministic control layer.

The **AI reasoning layer** handles the one task that benefits from language understanding: selecting approved research tools within a bounded ReAct loop and synthesising collected findings into a human-readable summary. Structured upgrade requests, risk factors, PR descriptions, and all safety-gate outcomes are produced by deterministic Python components.

The **deterministic control layer** handles all pass/fail decisions: Helm linting, template rendering, Kubernetes version comparison, health metric evaluation against thresholds, and risk score calculation. No LLM output directly determines whether a quality gate passes. This separation ensures that the system's safety properties are independent of LLM behaviour—the LLM can recommend, Python decides.

The state machine (Requested → Analysed → Planned → Validated → INTDeployed → AwaitingApproval → Completed; terminal failure states: Blocked, INTFailed, Paused) enforces workflow progression through deterministic gate evaluations only.

Key technical decisions:

- **Pydantic validation for all inter-agent messages:** strict schemas with `extra="forbid"` prevent schema drift between components and ensure LLM-generated content conforms to expected structure before influencing the workflow (Pydantic, n.d.).
- **GateResult.UNKNOWN for missing health evidence:** a missing monitoring snapshot never defaults to PASS. This conservative design prevents silent failures from masquerading as successes.
- **RiskLevel.UNKNOWN for incomplete research:** when the ReAct agent exhausts its iteration limit before collecting mandatory evidence, the risk cannot be assessed. The report records UNKNOWN rather than inheriting a low numeric score, and the coordinator transitions to Paused before any quality gate is evaluated.
- **Secret detection in rendered manifests:** all helm template output is scanned for credential patterns before being stored or passed to downstream components.
- **Approved-tool enforcement in the ReAct loop:** only the three registered retrieval functions may be called. Tool arguments are validated against strict Pydantic models and scope-checked against the active upgrade request. Any call to an unregistered tool, wrong component, or mismatched chart version raises an exception before execution, reducing the impact of prompt-injection attempts that try to redirect tool calls outside the active upgrade scope.

4. Technical Design Decisions and Tradeoffs

The most significant tradeoff in this project is between automation depth and safety. A fully automated system—one that could directly commit to GitOps repositories and trigger

production deployments—would be faster but would expose the platform to the consequences of AI errors at production scale. The design deliberately stops short: the system proposes changes to an outputs/ directory and flags `AwaitingApproval`; actual production deployment requires a human engineer to review and approve (Amodei et al., 2016).

A second tradeoff concerns the deterministic risk score versus a trained ML classifier. A model trained on upgrade history might produce more accurate risk predictions, but would be harder to explain and audit. In a platform engineering context, the ability to explain why a risk score is 65 rather than 35 is operationally essential.

A third tradeoff is in document retrieval. A vector embedding search (as used in the agentic-ai-udaplay RAG system) could improve recall from release notes, particularly for large or poorly structured documents. The current keyword-based retrieval is more predictable and auditable; embedding-based retrieval is identified as a future extension.

5. Ethical Considerations and Responsible AI

The principal ethical concerns in this domain are not demographic fairness (the typical focus in consumer AI systems) but rather automation risk to critical infrastructure.

Automation bias is the primary risk: an engineer who trusts AI recommendations uncritically may approve a production deployment without independently reviewing the evidence. The system mitigates this through mandatory human approval for all production deployments (enforced as a hard constraint), complete audit logs, and explicit risk level labels that require engineers to acknowledge the risk before acting.

Hallucinated compatibility claims are a concern in any LLM-assisted system. The design addresses this by treating all LLM output as hypotheses. LLM-generated risks are labelled as INFO-level, do not set safety-critical fields, and cannot directly determine compatibility or gate outcomes. Compatibility decisions are derived from registered retrieval tools and deterministic version checks.

Prompt injection through documentation is an explicit threat model: a malicious actor could embed instructions in a release note document. The system mitigates this risk by treating retrieved documents as data strings parsed for structured information by deterministic code; the LLM receives already-extracted findings, not raw document content.

Secret exposure is mitigated by scanning all Helm template output for credential patterns before storing or sharing results. Agents cannot access Kubernetes secrets, execute arbitrary shell commands, or write directly to production repositories.

Accountability remains with the platform engineer throughout. The AI system is a decision-support and workflow-coordination tool; it does not own the production service. Audit logs record decision summaries, tool actions, observations, state transitions, and outcomes without capturing private chain-of-thought, ensuring that every AI action is attributable and reviewable (European Commission, 2021).

6. Evaluation and Reflection

Five upgrade scenarios were executed through the integrated system:

Successful patch upgrade (kafka-connect 0.18.2 → 0.18.3): Risk score 20/100 (Low). All validation and health gates passed; system recommended promotion with human approval required.

Breaking value change (kafka-connect → 0.19.0, CP 8.0.2): Risk score 45/100 (Medium). The research agent detected the renamed configuration key from release notes; validation confirmed the deprecated value in existing environment files. All gates passed; system transitioned to AwaitingApproval, requiring human review before PROD promotion.

Kubernetes incompatibility (kafka-connect 0.19.0 on K8s 1.29 clusters): Risk score 70/100 (High). Compatibility check blocked the upgrade before any deployment; no GitOps changes were generated.

Failed INT rollout (kafka-connect 0.19.0, simulated crash): Risk score 45/100 (Medium). Health evaluation detected pod_ready_percent=50% and restart_count=8, both exceeding thresholds. System transitioned to INTFailed and recommended rollback.

Missing observability (prometheus 25.2.0, no metrics): Risk score 20/100 (Low), but health evaluation returned all-UNKNOWN gate results. System transitioned to Paused, demonstrating that UNKNOWN evidence is never treated as PASS.

The evaluation confirmed that all five scenarios produced the expected outcomes. The deterministic gate architecture successfully prevented incorrect promotions in all failure and uncertainty cases. The LLM synthesis step added contextual explanation to findings but was not required for any safety-critical decision.

A key limitation is that the system operates on simulated data. A production deployment would require integration with real Argo CD, real Prometheus metrics, and actual Helm chart repositories — none of which are appropriate for an academic prototype.

7. Professional Relevance

This project demonstrates capabilities directly relevant to senior platform engineering and AI engineering roles. The integration of statistical analysis, ML-inspired risk modelling, generative AI synthesis, and bounded agentic research into a single production-adjacent system reflects the cross-domain fluency that senior engineers bring to complex operational problems. The explicit separation of AI reasoning and deterministic control is particularly relevant in regulated industries where AI decisions must be explainable and auditable.

8. Future Extensions and Conclusion

Several extensions would strengthen the system for production use:

- **Vector embedding search** for release notes using ChromaDB (as used in agentic-ai-udaplay), improving recall for large or poorly structured documents.
- **Real Argo CD integration** via the Argo CD REST API to retrieve actual sync and health status rather than simulated snapshots.
- **Automated PR creation** in a safe test repository using the GitHub API, making the GitOps proposal end-to-end.
- **Risk model training** using historical upgrade data to replace the manually weighted rule set with a calibrated classifier.
- **Soak period monitoring** with real-time Prometheus query integration to observe deployment health over a 30-minute window.

This project brings together the analytical, modelling, generative, and agentic skills developed across the AI Mastery Capstone into a single, cohesive system. The design choices—deterministic safety gates, human approval requirements, evidence-linked recommendations, and complete audit trails—reflect the professional judgment required to deploy AI systems responsibly in critical infrastructure contexts.

References

- Amodei, D., Olah, C., Steinhardt, J., Christiano, P., Schulman, J., & Mané, D. (2016). *Concrete problems in AI safety*. arXiv preprint arXiv:1606.06565. <https://arxiv.org/abs/1606.06565>
- European Commission. (2021). *Proposal for a regulation of the European Parliament and of the Council laying down harmonised rules on Artificial Intelligence (Artificial Intelligence Act)*. COM/2021/206 final. This proposal formed the regulatory framework considered during system design; the final EU AI Act was adopted in 2024. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52021PC0206>
- Pydantic. (n.d.). *Pydantic documentation: Data validation using Python type annotations*. <https://docs.pydantic.dev/latest/>

- Sculley, D., Holt, G., Golovin, D., Davydov, E., Phillips, T., Ebner, D., ... & Dennison, D. (2015). Hidden technical debt in machine learning systems. *Advances in Neural Information Processing Systems*, 28. <https://proceedings.neurips.cc/paper/2015/hash/86df7dcfd896fcaf2674f757a2463eba-Abstract.html>
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *International Conference on Learning Representations (ICLR 2023)*. <https://arxiv.org/abs/2210.03629>